

Descrição do Problema e Primeira Agregação

January 5, 2021

```
[36]: import pandas as pd
import seaborn as sn

import base_functions

pd.set_option('display.float_format', lambda x: '%.3f' % x)
```

1 Desafio

Dentre os desafios propostos, o escolhido foi:

- *"Como prever o desempenho dos estudantes nas avaliações finais".*

Cada estudante em cada módulo faz em média quatro provas, e a nota final consiste de uma média ponderada dessas notas, que define a aprovação ou não do estudante. Destarte, pode-se pensar em "avaliações finais" em dois níveis de agregação distintos: - primeiro é o de tentar prever o desempenho em cada prova, uma vez que não há uma avaliação final *per se*, e sim um conjunto de avaliações para cada módulo; - a outra interpretação possível é entender a avaliação final como a média das avaliações que é também o resultado final de desempenho.

Aqui utilizamos a segunda interpretação, por motivo de simplicidade e performance preditiva, uma vez que ambas foram testadas. Os testes e análises não foram exaustivamente incluídos nesse texto visando manter uma linha condutória com leitura mais fluída.

Pontuo também que o texto no formato markdown será feito em português, e os comentários relacionados as operações computacionais e feitos na forma de comentários no código serão feitos em inglês por convenção.

2 Entendimento dos dados

Essa seção consiste de três etapas: 1. Apresentação da descrição dos dados presente no link https://analyse.kmi.open.ac.uk/open_dataset.

2. Exposição dos dados e de algumas propriedades *e.g* cardinalidade.
3. Agregação dos diferentes datasets visando criar um dataset analítico inicial.

O banco de dados é organizado em sete diferentes tabelas, descritas na imagem acima. Descrevendo de forma simples, as tabelas são divididas nas que descrevem: - O estudante - Os módulos - As atividades no VLE (virtual learning environment).

O objetivo de agregação é criar uma tabela única que tenha como chave composta a triade [**Estudante** | **Curso** | **Semestre**]. Dessa forma, assume-se que o mesmo estudante, cursando cursos diferentes, ou o mesmo curso em outro momento do tempo será tratado de forma independente.

3 Descrição do problema, variáveis, e visualização dos dados de cada tabela

As descrições foram tiradas integralmente da página de referência.

Description "This page introduces the anonymised Open University Learning Analytics Dataset (OULAD). It contains data about courses, students and their interactions with Virtual Learning Environment (VLE) for seven selected courses (called modules). Presentations of courses start in February and October - they are marked by "B" and "J" respectively. The dataset consists of tables connected using unique identifiers. All tables are stored in the csv format."

```
[21]: df_0 = pd.read_csv("../data/assessments.csv")
```

This file contains information about assessments in module-presentations. Usually, every presentation has a number of assessments followed by the final exam.

code_module - identification code of the module, to which the assessment belongs.
code_presentation - identification code of the presentation, to which the assessment belongs.
id_assessment - identification number of the assessment. **assessment_type** - type of assessment. Three types of assessments exist: Tutor Marked Assessment (TMA), Computer Marked Assessment (CMA) and Final Exam (Exam). **date** - information about the final submission date of the assessment calculated as the number of days since the start of the module-presentation. The starting date of the presentation has number 0 (zero). **weight** - weight of the assessment in %. Typically, Exams are treated separately and have the weight 100%; the sum of all other assessments is 100%.

```
[22]: df_0.head(3)
```

```
[22]:
```

	code_module	code_presentation	id_assessment	assessment_type	date	weight
0	AAA	2013J	1752	TMA	19.0	10.0
1	AAA	2013J	1753	TMA	54.0	20.0
2	AAA	2013J	1754	TMA	117.0	20.0

```
[8]: base_functions.describe_(df_0)
```

```
shape (206, 6)
#####
Contagem de categorias: code_module
FFF      52
BBB      42
```

```

DDD      35
GGG      30
CCC      20
EEE      15
AAA      12
Name: code_module, dtype: int64
#####
Contagem de categorias: code_presentation
2014J     57
2014B     57
2013J     53
2013B     39
Name: code_presentation, dtype: int64
#####
Contagem de categorias: assessment_type
TMA       106
CMA       76
Exam      24
Name: assessment_type, dtype: int64

```

```
[10]: df_1 = pd.read_csv("../data/courses.csv")
```

File contains the list of all available modules and their presentations. **code_module** – code name of the module, which serves as the identifier. **code_presentation** – code name of the presentation. It consists of the year and “B” for the presentation starting in February and “J” for the presentation starting in October. **length** - length of the module-presentation in days. The structure of B and J presentations may differ and therefore it is good practice to analyse the B and J presentations separately. Nevertheless, for some presentations the corresponding previous B/J presentation do not exist and therefore the J presentation must be used to inform the B presentation or vice versa. In the dataset this is the case of CCC, EEE and GGG modules.

```
[7]: df_1.head(3)
```

```
[7]:
```

	code_module	code_presentation	module_presentation_length
0	AAA	2013J	268
1	AAA	2014J	269
2	BBB	2013J	268

```
[11]: base_functions.describe_(df_1)
```

```

shape (22, 3)
#####
Contagem de categorias: code_module
BBB      4
FFF      4
DDD      4
EEE      3
GGG      3

```

```

AAA      2
CCC      2
Name: code_module, dtype: int64
#####
Contagem de categorias: code_presentation
2014J      7
2014B      6
2013J      6
2013B      3
Name: code_presentation, dtype: int64

```

```
[14]: # join df_0 and df_1
df_0_1 = df_0.merge(df_1, on=['code_module', 'code_presentation'])
```

```
[12]: df_2 = pd.read_csv("../data/studentAssessment.csv")
```

This file contains the results of students' assessments.

If the student does not submit the assessment, no result is recorded. The final exam submissions is missing, if the result of the assessments is not stored in the system. This file contains the following columns: **id_assessment** – the identification number of the assessment. **id_student** – a unique identification number for the student. **date_submitted** – the date of student submission, measured as the number of days since the start of the module presentation. **is_banked** – a status flag indicating that the assessment result has been transferred from a previous presentation. **score** – the student's score in this assessment. The range is from 0 to 100. The score lower than 40 is interpreted as Fail. The marks are in the range from 0 to 100.

```
[11]: df_2.head(3)
```

```
[11]:
```

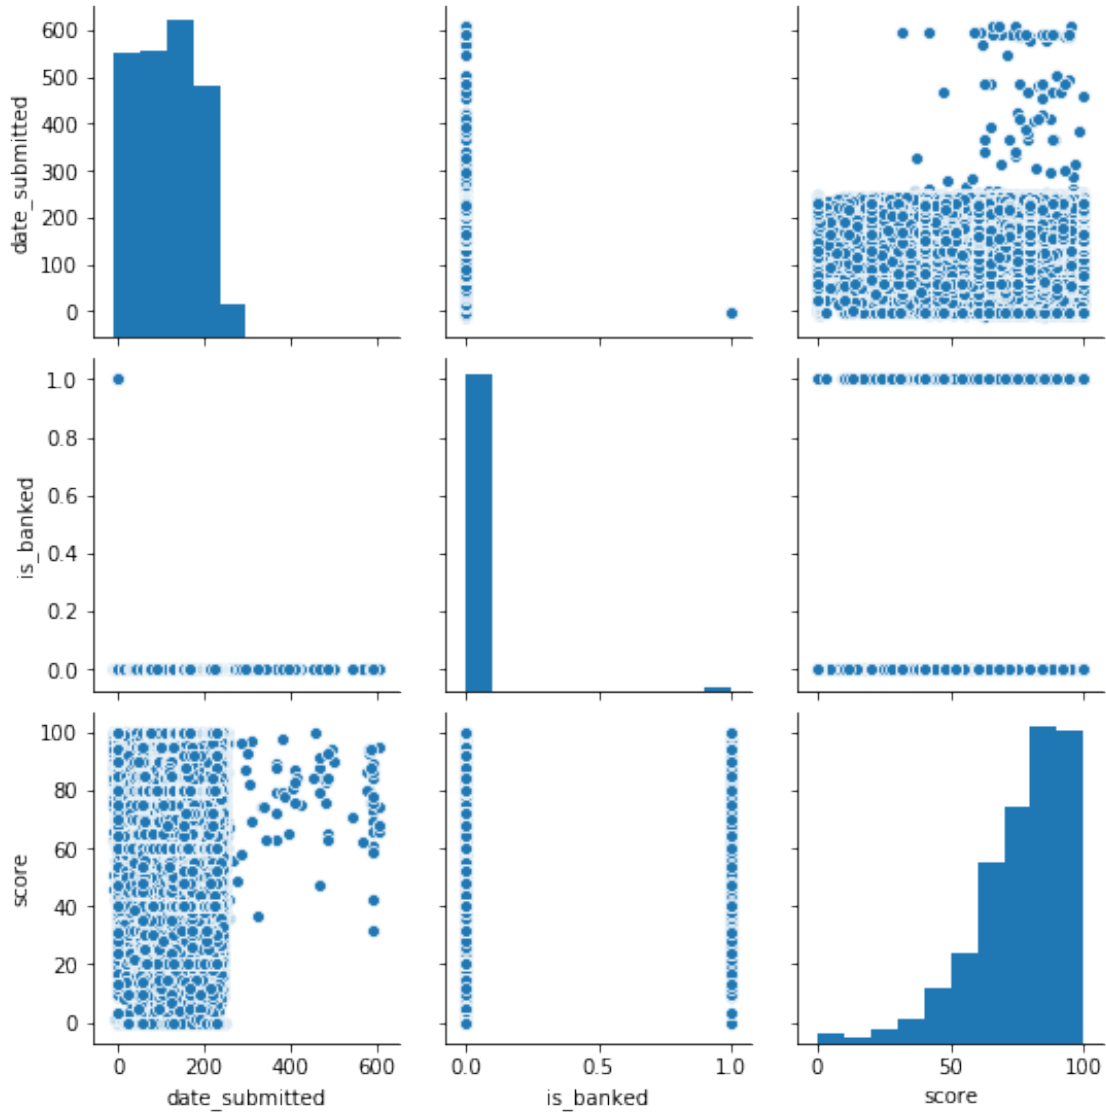
	id_assessment	id_student	date_submitted	is_banked	score
0	1752	11391	18	0	78.0
1	1752	28400	22	0	70.0
2	1752	31604	17	0	72.0

```
[13]: df_2.shape
```

```
[13]: (173912, 5)
```

```
[26]: sn.pairplot(df_2.drop(columns=['id_assessment', 'id_student']))
```

```
[26]: <seaborn.axisgrid.PairGrid at 0x7f2906657898>
```



```
[49]: # join df_2 with the previous two.
df_0_1_2 = df_2.merge(df_0_1, on = ['id_assessment'])
```

```
[48]: df_3 = pd.read_csv("../data/studentInfo.csv")
```

This file contains demographic information about the students together with their results. **code_module** – an identification code for a module on which the student is registered. **code_presentation** - the identification code of the presentation during which the student is registered on the module. **id_student** – a unique identification number for the student. **gender** – the student’s gender. **region** – identifies the geographic region, where the student lived while taking the module-presentation. **highest_education** – highest student education level on entry to the module presentation. **imd_band** – specifies the Index of Multiple Deprivation band of the place where the student lived during the module-presentation. **age_band** – band of the student’s

age. **num_of_prev_attempts** – the number times the student has attempted this module. **studied_credits** – the total number of credits for the modules the student is currently studying. **disability** – indicates whether the student has declared a disability. **final_result** – student's final result in the module-presentation.

```
[13]: df_3.head(3)
```

```
[13]:  code_module code_presentation  id_student gender      region \
0      AAA      2013J      11391      M  East Anglian Region
1      AAA      2013J      28400      F      Scotland
2      AAA      2013J      30268      F  North Western Region

      highest_education imd_band age_band  num_of_prev_attempts \
0      HE Qualification  90-100%  55<=      0
1      HE Qualification  20-30%  35-55      0
2  A Level or Equivalent  30-40%  35-55      0

      studied_credits disability final_result
0      240      N      Pass
1      60      N      Pass
2      60      Y  Withdrawn
```

```
[37]: desc(df_3)
```

```
shape (32593, 12)
code_module
BBB      7909
FFF      7762
DDD      6272
CCC      4434
EEE      2934
GGG      2534
AAA       748
Name: code_module, dtype: int64
code_presentation
2014J      11260
2013J      8845
2014B       7804
2013B       4684
Name: code_presentation, dtype: int64
gender
M      17875
F      14718
Name: gender, dtype: int64
region
Scotland      3446
East Anglian Region  3340
London Region    3216
```

```

South Region          3092
North Western Region  2906
West Midlands Region  2582
South West Region     2436
East Midlands Region  2365
South East Region     2111
Wales                 2086
Yorkshire Region      2006
North Region          1823
Ireland               1184
Name: region, dtype: int64
highest_education
A Level or Equivalent 14045
Lower Than A Level    13158
HE Qualification       4730
No Formal quals        347
Post Graduate Qualification 313
Name: highest_education, dtype: int64
imd_band
20-30%    3654
30-40%    3539
10-20     3516
0-10%     3311
40-50%    3256
50-60%    3124
60-70%    2905
70-80%    2879
80-90%    2762
90-100%   2536
Name: imd_band, dtype: int64
age_band
0-35      22944
35-55      9433
55<=       216
Name: age_band, dtype: int64
disability
N      29429
Y       3164
Name: disability, dtype: int64
final_result
Pass      12361
Withdrawn 10156
Fail       7052
Distinction 3024
Name: final_result, dtype: int64

```

[]:

```
[50]: # join df_3 with the previous three.
df_0_1_2_3 = df_0_1_2.merge(df_3, on=['code_module', 'code_presentation',
↳ 'id_student'])
```

```
[ ]:
```

```
[51]: df_4 = pd.read_csv("../data/studentRegistration.csv")
```

This file contains information about the time when the student registered for the module presentation. For students who unregistered the date of unregistration is also recorded. **code_module** – an identification code for a module. **code_presentation** – the identification code of the presentation. **id_student** – a unique identification number for the student. **date_registration** – the date of student's registration on the module presentation, this is the number of days measured relative to the start of the module-presentation (e.g. the negative value -30 means that the student registered to module presentation 30 days before it started). **date_unregistration** – date of student unregistration from the module presentation, this is the number of days measured relative to the start of the module-presentation. Students, who completed the course have this field empty. Students who unregistered have Withdrawal as the value of the final_result column in the studentInfo.csv file.


```
[15]: df_4.head(3)
```

```
[15]:   code_module code_presentation  id_student  date_registration \
0          AAA          2013J         11391          -159.0
1          AAA          2013J         28400           -53.0
2          AAA          2013J         30268           -92.0

      date_unregistration
0                   NaN
1                   NaN
2                  12.0
```

```
[29]: sn.pairplot(df_4.drop(columns=['id_student']))
```

```
[29]: <seaborn.axisgrid.PairGrid at 0x7f28fe8ae128>
```

```
[ ]:
```

```
[52]: # join df_4 with the previous four.
df_0_1_2_3_4 = df_0_1_2_3.merge(df_4, on=['code_module', 'code_presentation',
↪ 'id_student'])
```

```
[53]: df_5 = pd.read_csv("../data/studentVle.csv")
```


The `studentVle.csv` file contains information about each student's interactions with the materials in the VLE. (virtual learning environment) **code_module** – an identification code for a module. **code_presentation** – the identification code of the module presentation. **id_student** – a unique identification number for the student. **id_site** – an identification number for the VLE material. **date** – the date of student's interaction with the material measured as the number of days since the start of the module-presentation. **sum_click** – the number of times a student interacts with the material in that day.

```
[17]: df_5.head(3)
```

```
[17]: code_module code_presentation id_student id_site date sum_click
0      AAA      2013J      28400    546652   -10         4
1      AAA      2013J      28400    546652   -10         1
2      AAA      2013J      28400    546652   -10         1
```

```
[32]: sn.pairplot(df_5.drop(columns=['id_site', 'id_student']))
```

```
[32]: <seaborn.axisgrid.PairGrid at 0x7f28fe9c6e10>
```



```
[37]: df_5['sum_click'].describe()
```

```
[37]: count    10655280.000
      mean         3.717
      std         8.849
      min         1.000
      25%         1.000
      50%         2.000
```

```
75%          3.000
max         6977.000
Name: sum_click, dtype: float64
```

```
[54]: # Here we separate just the sum_click variable, which is the one that keeps a
# meaningful interpretation after aggregation. We could easily perform this
# aggregation on the
# next step of the analysis instead of now.
grouped_df_5 = df_5.drop(columns=[
    'id_site', 'date']).groupby(
    by=['code_module', 'code_presentation', 'id_student'])
    .sum().reset_index()

# Here we aggregate the df_5 with all the previous.
df_0_1_2_3_4_5 = df_0_1_2_3_4.merge(
    grouped_df_5, on=['code_module', 'code_presentation', 'id_student'])
```

```
[18]: df_6 = pd.read_csv("../data/vle.csv")
```

The csv file contains information about the available materials in the VLE. Typically these are html pages, pdf files, etc. Students have access to these materials online and their interactions with the materials are recorded. **id_site** – an identification number of the material. **code_module** – an identification code for module. **code_presentation** - the identification code of presentation. **activity_type** – the role associated with the module material. **week_from** – the week from which the material is planned to be used. **week_to** – week until which the material is planned to be used.

```
[19]: df_6.head(3)
```

```
[19]:   id_site code_module code_presentation activity_type week_from week_to
0    546943         AAA         2013J         resource         NaN         NaN
1    546712         AAA         2013J         oucontent         NaN         NaN
2    546998         AAA         2013J         resource         NaN         NaN
```

```
[76]: df_6.shape
```

```
[76]: (6364, 6)
```

```
[15]: # I chose to not include the df_6 in the analytical df, since I understand that
# the activity
# type var couldn't be grouped by student, and the week_from and week_to have
# no simple
# interpretation once we group the data for each student.
```

3.0.1 Exportação do primeiro dataset analítico.

```
[98]: df_0_1_2_3_4_5.to_csv('base_data.csv', index=False)
```